

```

1 import pandas as pd
2 import re
3 import torch
4 import numpy as np
5 import matplotlib.pyplot as plt
6
7 from sklearn.model_selection import train_test_split
8 from sklearn.metrics import accuracy_score, classification_report, confusion_matrix, ConfusionMatrixDisplay
9
10 from transformers import (
11     AutoTokenizer,
12     AutoModelForSequenceClassification,
13     Trainer,
14     TrainingArguments
15 )
16
17 from torch.utils.data import Dataset
18

```

```

1 df = pd.read_csv("BTC_tweets_daily_example.csv")
2 print(df.head())
3

Unnamed: 0 Date \
0 0 Fri Mar 23 00:40:32 +0000 2018
1 1 Fri Mar 23 00:40:34 +0000 2018
2 2 Fri Mar 23 00:40:35 +0000 2018
3 3 Fri Mar 23 00:40:36 +0000 2018
4 4 Fri Mar 23 00:40:36 +0000 2018

Tweet Screen_name \
0 RT @ALXTOKEN: Paul Krugman, Nobel Luddite. I h... myresumerocket
1 @lopp @Kevin_Pham @psycho_sage @naval But @Pr... BitMacro
2 RT @tippereconomy: Another use case for #block... hojachotpur
3 free coins https://t.co/DiuoePJdap denies_distro
4 RT @payvxofficial: WE are happy to announce th... aditzgraha

Source \
0 []
1 [u'Bitcoin']
2 [u'blockchain', u'Tipper', u'TipperEconomy']
3 []
4 []

Link Sentiment \
0 <a href="http://twitter.com" rel="nofollow">Tw... ['neutral']
1 <a href="http://twitter.com/download/android" ... ['neutral']
2 <a href="http://twitter.com" rel="nofollow">Tw... ['positive']
3 <a href="http://twitter.com" rel="nofollow">Tw... ['positive']
4 <a href="http://twitter.com/download/android" ... ['positive']

sent_score New_Sentiment_Score New_Sentiment_State
0 0.0 0.00000 0.0
1 0.0 0.00000 0.0
2 1.0 0.136364 1.0
3 1.0 0.400000 1.0
4 1.0 0.468182 1.0

```

```

1 def clean_text(text):
2     text = str(text).lower()
3     text = re.sub(r"http\S+", "", text)
4     text = re.sub(r"@\\w+", "", text)
5     text = re.sub(r"#\\w+", "", text)
6     text = re.sub(r"[^a-z\\s]", "", text)
7     return text
8

```

```

1 df["Tweet"] = df["Tweet"].astype(str)
2 df["Tweet"] = df["Tweet"].apply(clean_text)
3
4 # Remove rows without labels
5 df = df.dropna(subset=["New_Sentiment_State"])
6

```

```

1 label_mapping = {
2     label: idx for idx, label in enumerate(sorted(df["New_Sentiment_State"].unique()))
3 }
4
5 df["label"] = df["New_Sentiment_State"].map(label_mapping)
6

```

```

7 num_labels = len(label_mapping)
8 print("Label mapping:", label_mapping)
9

```

```
Label mapping: {np.float64(-1.0): 0, np.float64(0.0): 1, np.float64(1.0): 2}
```

```

1 X_train, X_test, y_train, y_test = train_test_split(
2     df["Tweet"],
3     df["label"],
4     test_size=0.2,
5     random_state=42,
6     stratify=df["label"]
7 )
8

```

```

1 MODEL_NAME = "distilroberta-base"
2
3 tokenizer = AutoTokenizer.from_pretrained(MODEL_NAME)
4
5 device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
6
7 model = AutoModelForSequenceClassification.from_pretrained(
8     MODEL_NAME,
9     num_labels=num_labels
10 ).to(device)

```

```
/usr/local/lib/python3.12/dist-packages/huggingface_hub/utils/_auth.py:94: UserWarning:
The secret `HF_TOKEN` does not exist in your Colab secrets.
To authenticate with the Hugging Face Hub, create a token in your settings tab (https://huggingface.co/settings/tokens), set
You will be able to reuse this secret in all of your notebooks.
Please note that authentication is recommended but still optional to access public models or datasets.
    warnings.warn(
Loading weights: 100%                                         101/101 [00:00<00:00, 314.86it/s, Materializing param=roberta.encoder.layer.5.output.dense.weight]
RobertaForSequenceClassification LOAD REPORT from: distilroberta-base
Key           | Status   |
-----+-----+
lm_head.layer_norm.bias | UNEXPECTED |
roberta.pooler.dense.weight | UNEXPECTED |
lm_head.dense.weight | UNEXPECTED |
lm_head.layer_norm.weight | UNEXPECTED |
roberta.pooler.dense.bias | UNEXPECTED |
lm_head.dense.bias | UNEXPECTED |
lm_head.bias | UNEXPECTED |
classifier.dense.weight | MISSING |
classifier.out_proj.weight | MISSING |
classifier.out_proj.bias | MISSING |
classifier.dense.bias | MISSING |

```

Notes:

- UNEXPECTED :can be ignored when Loading from different task/architecture; not ok if you expect identical arch.
- MISSING :those params were newly initialized because missing from the checkpoint. Consider training on your downstream task.

```

1 class TweetDataset(Dataset):
2     def __init__(self, texts, labels):
3         self.encodings = tokenizer(
4             texts.tolist(),
5             truncation=True,
6             padding=True,
7             max_length=128
8         )
9         self.labels = labels.tolist()
10
11    def __getitem__(self, idx):
12        item = {k: torch.tensor(v[idx]) for k, v in self.encodings.items()}
13        item["labels"] = torch.tensor(self.labels[idx])
14        return item
15
16    def __len__(self):
17        return len(self.labels)
18

```

```

1 train_dataset = TweetDataset(X_train, y_train)
2 test_dataset = TweetDataset(X_test, y_test)
3

```

```

1 training_args = TrainingArguments(
2     output_dir="./results",
3     per_device_train_batch_size=32,
4     per_device_eval_batch_size=32,

```

```

5     num_train_epochs=1,           # 🔥 FAST
6     learning_rate=2e-5,
7     weight_decay=0.01,
8     logging_steps=100,
9     save_strategy="no",
10    eval_strategy="no",
11    fp16=True,                  # 🔥 MUCH FASTER ON GPU
12    report_to="none"
13 )

```

```

1 trainer = Trainer(
2     model=model,
3     args=training_args,
4     train_dataset=train_dataset
5 )

```

```
1 trainer.train()
```

[1272/1272 01:47, Epoch 1/1]

Step Training Loss

| | |
|------|----------|
| 100 | 0.824642 |
| 200 | 0.438055 |
| 300 | 0.342024 |
| 400 | 0.292431 |
| 500 | 0.248264 |
| 600 | 0.246954 |
| 700 | 0.171308 |
| 800 | 0.195838 |
| 900 | 0.157617 |
| 1000 | 0.153829 |
| 1100 | 0.161120 |
| 1200 | 0.164257 |

```
TrainOutput(global_step=1272, training_loss=0.2765104238342189, metrics={'train_runtime': 109.0079, 'train_samples_per_second': 373.193, 'train_steps_per_second': 11.669, 'total_flos': 831506221697598.0, 'train_loss':
```

```

1 predictions = trainer.predict(test_dataset)
2 y_pred = np.argmax(predictions.predictions, axis=1)
3
4 print("Accuracy:", accuracy_score(y_test, y_pred))
5 print("\nClassification Report:\n")
6 print(classification_report(y_test, y_pred, target_names=[str(key) for key in label_mapping.keys()]))

```

Accuracy: 0.9584111690099302


Classification Report:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| -1.0 | 0.93 | 0.87 | 0.90 | 1313 |
| 0.0 | 0.95 | 0.98 | 0.96 | 4349 |
| 1.0 | 0.97 | 0.97 | 0.97 | 4509 |
| accuracy | | | 0.96 | 10171 |
| macro avg | 0.95 | 0.94 | 0.94 | 10171 |
| weighted avg | 0.96 | 0.96 | 0.96 | 10171 |

```

1 cm = confusion_matrix(y_test, y_pred)
2
3 disp = ConfusionMatrixDisplay(confusion_matrix=cm)
4 disp.plot(cmap="Blues")
5 plt.title("DistilRoBERTa Confusion Matrix")
6 plt.show()
7

```




```
1 Start coding or generate with AI.
```

```

1 import pandas as pd
2 import matplotlib.pyplot as plt
3
4 # Convert numeric labels to Series
5 actual_counts = pd.Series(y_test).value_counts().sort_index()
6 pred_counts = pd.Series(y_pred).value_counts().sort_index()
7
8 # Align indexes so missing classes don't cause NaN
9 df_plot = pd.DataFrame({
10     "Actual": actual_counts,
11     "Predicted": pred_counts
12 }).fillna(0)
13
14 # Plot
15 df_plot.plot(
16     kind="bar",
17     figsize=(6,4)
18 )
19
20 plt.title("Actual vs Predicted Sentiment Distribution")
21 plt.xlabel("Sentiment Class")
22 plt.ylabel("Count")
23 plt.xticks(rotation=0)
24 plt.tight_layout()
25 plt.show()
26

```




```


1 import matplotlib.pyplot as plt
2
3 neg_pos = {
4     "Negative": pred_counts.get(0, 0),    # class 0

```

```
5     "Positive": pred_counts.get(2, 0)    # class 2
6 }
7
8 plt.bar(neg_pos.keys(), neg_pos.values())
9 plt.title("Negative vs Positive Predictions")
10 plt.ylabel("Count")
11 plt.tight_layout()
12 plt.show()
13
```



```
1 acc = accuracy_score(y_test, y_pred)
2
3 plt.bar(["Accuracy"], [acc])
4 plt.ylim(0, 1)
5 plt.title("Model Accuracy")
6 plt.show()
7
```



```
1 import torch
2 from transformers import AutoTokenizer, AutoModelForSequenceClassification
3
4 # Load DistilRoBERTa
5 MODEL_NAME = "distilroberta-base"
6
7 tokenizer = AutoTokenizer.from_pretrained(MODEL_NAME)
8 model = AutoModelForSequenceClassification.from_pretrained(
9     MODEL_NAME,
10     num_labels=3
11 )
12
13 model.eval()
14
```

```

15 # Input sentence
16 sentence = "I just love how Bitcoin destroyed my savings today"
17
18 # Tokenize
19 inputs = tokenizer(
20     sentence,
21     return_tensors="pt",
22     truncation=True,
23     padding=True
24 )
25
26 # Predict
27 with torch.no_grad():
28     outputs = model(**inputs)
29     probs = torch.softmax(outputs.logits, dim=1)
30
31 # Labels (model output space: 0,1,2)
32 labels = {
33     0: "Negative 😞",
34     1: "Neutral 😐",
35     2: "Positive 😊"
36 }
37
38 # Max-probability decision
39 pred_id = torch.argmax(probs).item()
40
41 # Output
42 print("Sentence:", sentence)
43 print("Probabilities:")
44 print("Negative:", round(probs[0][0].item(), 3))
45 print("Neutral :", round(probs[0][1].item(), 3))
46 print("Positive:", round(probs[0][2].item(), 3))
47 print("\nFinal Prediction (max prob):", labels[pred_id])
48

```

Loading weights: 100% 101/101 [00:00<00:00, 414.06it/s, Materializing param=roberta.encoder.layer.5.output.dense.weight]

RobertaForSequenceClassification LOAD REPORT from: distilroberta-base

| Key | Status |
|-----------------------------|------------|
| lm_head.layer_norm.bias | UNEXPECTED |
| roberta.pooler.dense.weight | UNEXPECTED |
| lm_head.dense.weight | UNEXPECTED |
| lm_head.layer_norm.weight | UNEXPECTED |
| roberta.pooler.dense.bias | UNEXPECTED |
| lm_head.dense.bias | UNEXPECTED |
| lm_head.bias | UNEXPECTED |
| classifier.dense.weight | MISSING |
| classifier.out_proj.weight | MISSING |
| classifier.out_proj.bias | MISSING |
| classifier.dense.bias | MISSING |

Notes:

- UNEXPECTED :can be ignored when loading from different task/architecture; not ok if you expect identical arch.
- MISSING :those params were newly initialized because missing from the checkpoint. Consider training on your downstream task.

Sentence: I just love how Bitcoin destroyed my savings today

Probabilities:

Negative: 0.362
 Neutral : 0.324
 Positive: 0.315

Final Prediction (max prob): Negative 😞